

# Tooling Seminar - Titanic

May 4, 2017

```
In [2]: %matplotlib inline
        # Scientific computing
        import numpy as np

        # Data frames (R-style)
        import pandas as pd

        # Visualisation
        import matplotlib.pyplot as plt
        import seaborn as sns

        # Machine learning packages
        from sklearn import tree
        from sklearn.tree import export_graphviz
        from sklearn.metrics import confusion_matrix, f1_score, roc_auc_score
        from sklearn.ensemble import RandomForestClassifier
```

```
In [3]: # Read data
        train_data = pd.read_csv('./data/train.csv')
        test_data = pd.read_csv('./data/test.csv')
        test_y = pd.read_csv('./data/gender_submission.csv').Survived

        train_data.head(10)
```

```
Out[3]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
5	6	0	3	
6	7	0	1	
7	8	0	3	
8	9	1	3	
9	10	1	2	

```
Name    Sex  Age  SibSp  \
```

0		Braund, Mr. Owen Harris	male	22	1
1	Cummings, Mrs. John Bradley (Florence Briggs Th...		female	38	1
2		Heikkinen, Miss. Laina	female	26	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35	1
4		Allen, Mr. William Henry	male	35	0
5		Moran, Mr. James	male	NaN	0
6		McCarthy, Mr. Timothy J	male	54	0
7		Palsson, Master. Gosta Leonard	male	2	3
8	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)		female	27	0
9		Nasser, Mrs. Nicholas (Adele Achem)	female	14	1

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
5	0	330877	8.4583	NaN	Q
6	0	17463	51.8625	E46	S
7	1	349909	21.0750	NaN	S
8	2	347742	11.1333	NaN	S
9	0	237736	30.0708	NaN	C

```
In [4]: train_data.info()
print("-----")
test_data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 0 to 890
Data columns (total 12 columns):
PassengerId    891 non-null int64
Survived       891 non-null int64
Pclass         891 non-null int64
Name           891 non-null object
Sex            891 non-null object
Age           714 non-null float64
SibSp         891 non-null int64
Parch         891 non-null int64
Ticket        891 non-null object
Fare          891 non-null float64
Cabin         204 non-null object
Embarked      889 non-null object
dtypes: float64(2), int64(5), object(5)
memory usage: 90.5+ KB
-----

<class 'pandas.core.frame.DataFrame'>
Int64Index: 418 entries, 0 to 417
Data columns (total 11 columns):
```

```

PassengerId    418 non-null int64
Pclass         418 non-null int64
Name           418 non-null object
Sex            418 non-null object
Age           332 non-null float64
SibSp          418 non-null int64
Parch          418 non-null int64
Ticket         418 non-null object
Fare           417 non-null float64
Cabin          91 non-null object
Embarked       418 non-null object
dtypes: float64(2), int64(4), object(5)
memory usage: 39.2+ KB

```

As shown we have 891 entries in the training data and 418 entries in the test data. However, the info also shows there are missing values in the age column. For test there is also a missing Fare value.

```

In [5]: grouped = train_data.groupby('Survived')
        grouped.count()

```

```

Out[5]:
      PassengerId  Pclass  Name  Sex  Age  SibSp  Parch  Ticket  Fare  \
Survived
0             549     549   549  549  424    549    549    549    549
1             342     342   342  342  290    342    342    342    342

      Cabin  Embarked
Survived
0         68      549
1        136      340

```

From the breakdown above we see the our target class (Survived) is relatively balanced over the training set. We also again see missing values.

There are some columns (Name, Ticket, Embarked) that probably are not as strong for our prediction. They are also fairly difficult to transform into something more useful. Lets discard these.

After we get rid of the columns specified above, we still have to deal with the NaN values we found.

```

In [6]: train_data.drop(['PassengerId', 'Name', 'Ticket', 'Cabin', 'Embarked'], axis=1, inplace=True)
        train_data['Age'].fillna(train_data['Age'].mean(), inplace=True)
        train_data.Age = train_data.Age.round(decimals=0)
        train_data.head(10)

```

```

Out[6]:
   Survived  Pclass   Sex  Age  SibSp  Parch   Fare
0         0       3  male   22     1     0  7.2500
1         1       1 female   38     1     0 71.2833
2         1       3 female   26     0     0  7.9250

```

3	1	1	female	35	1	0	53.1000
4	0	3	male	35	0	0	8.0500
5	0	3	male	30	0	0	8.4583
6	0	1	male	54	0	0	51.8625
7	0	3	male	2	3	1	21.0750
8	1	3	female	27	0	2	11.1333
9	1	2	female	14	1	0	30.0708

```
In [7]: test_data.drop(['PassengerId', 'Name', 'Ticket', 'Cabin', 'Embarked'], axis=1, inplace=True)
test_data['Age'].fillna(test_data['Age'].mean(), inplace=True)
test_data.Age = test_data.Age.round(decimals=0)
test_data['Fare'].fillna(test_data['Fare'].median(), inplace=True)
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 418 entries, 0 to 417
Data columns (total 6 columns):
Pclass      418 non-null int64
Sex          418 non-null object
Age          418 non-null float64
SibSp        418 non-null int64
Parch        418 non-null int64
Fare         418 non-null float64
dtypes: float64(2), int64(3), object(1)
memory usage: 22.9+ KB
```

In our final preparation step we turn the 'Sex' column into a dummy column with 0 being male and 1 being female.

```
In [8]: #train_data.Pclass = train_data.Pclass.astype('category')
#train_data.Sex = train_data.Sex.astype('category')
train_data['Sex'] = np.where(train_data['Sex'] == 'female', 1, 0)

#test_data.Pclass = test_data.Pclass.astype('category')
#test_data.Sex = test_data.Sex.astype('category')
test_data['Sex'] = np.where(test_data['Sex'] == 'female', 1, 0)
test_data.head(10)
```

```
Out [8]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare
0	3	0	34	0	0	7.8292
1	3	1	47	1	0	7.0000
2	2	0	62	0	0	9.6875
3	3	0	27	0	0	8.6625
4	3	1	22	1	1	12.2875
5	3	0	14	0	0	9.2250
6	3	1	30	0	0	7.6292
7	2	0	26	1	1	29.0000
8	3	1	18	0	0	7.2292
9	3	0	21	2	0	24.1500

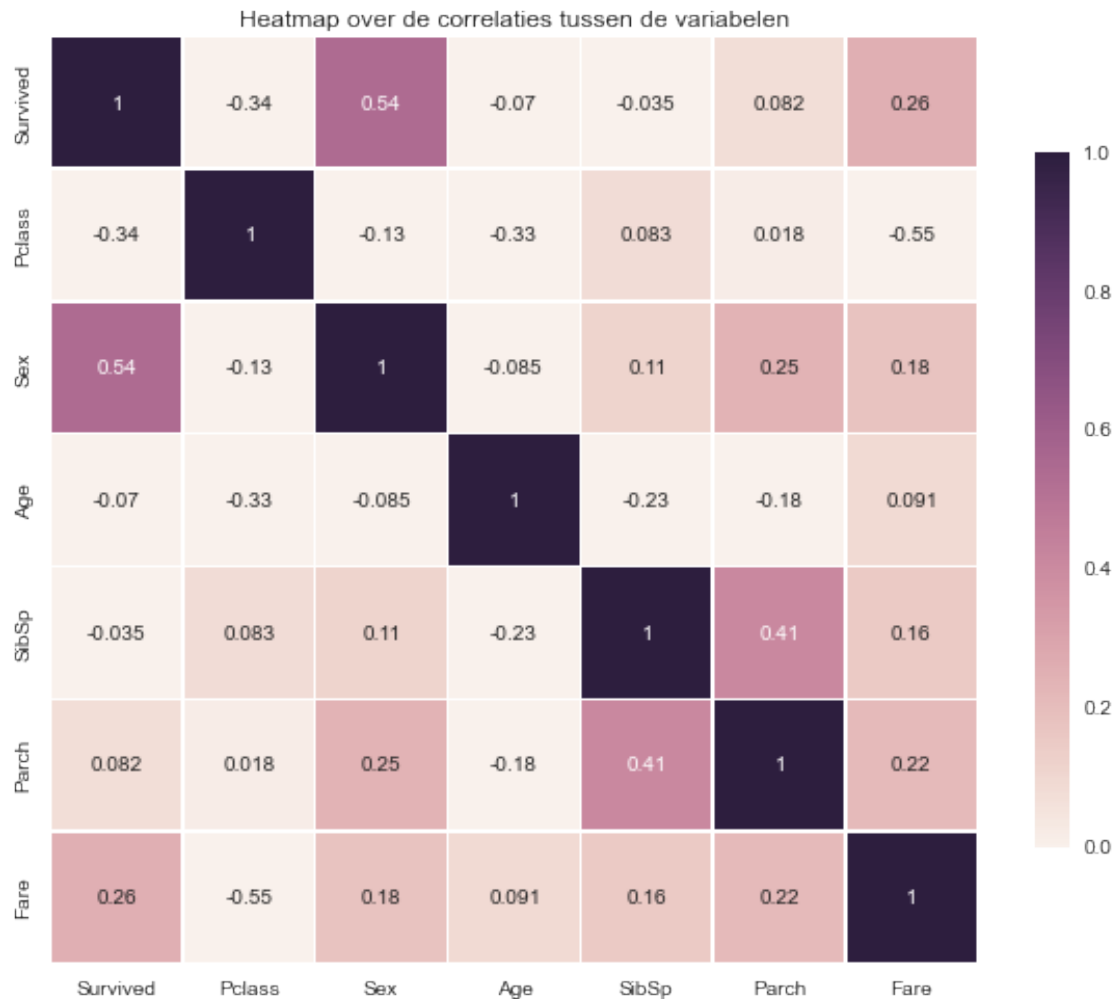
Lets start looking at our data now we have cleaned it. Do we need any other steps to handle correlations before we train a model?

```
In [9]: correlations = train_data.corr()
        print(correlations['Survived'])

        plt.figure(figsize=(10,10))
        sns.heatmap(correlations,linewidths=0.25, square=True, cbar_kws={'shrink' : .6}, annot=
        plt.title("Heatmap over de correlaties tussen de variabelen")
```

```
Survived    1.000000
Pclass      -0.338481
Sex          0.543351
Age         -0.070324
SibSp       -0.035322
Parch        0.081629
Fare         0.257307
Name: Survived, dtype: float64
```

```
Out[9]: <matplotlib.text.Text at 0x938f550>
```

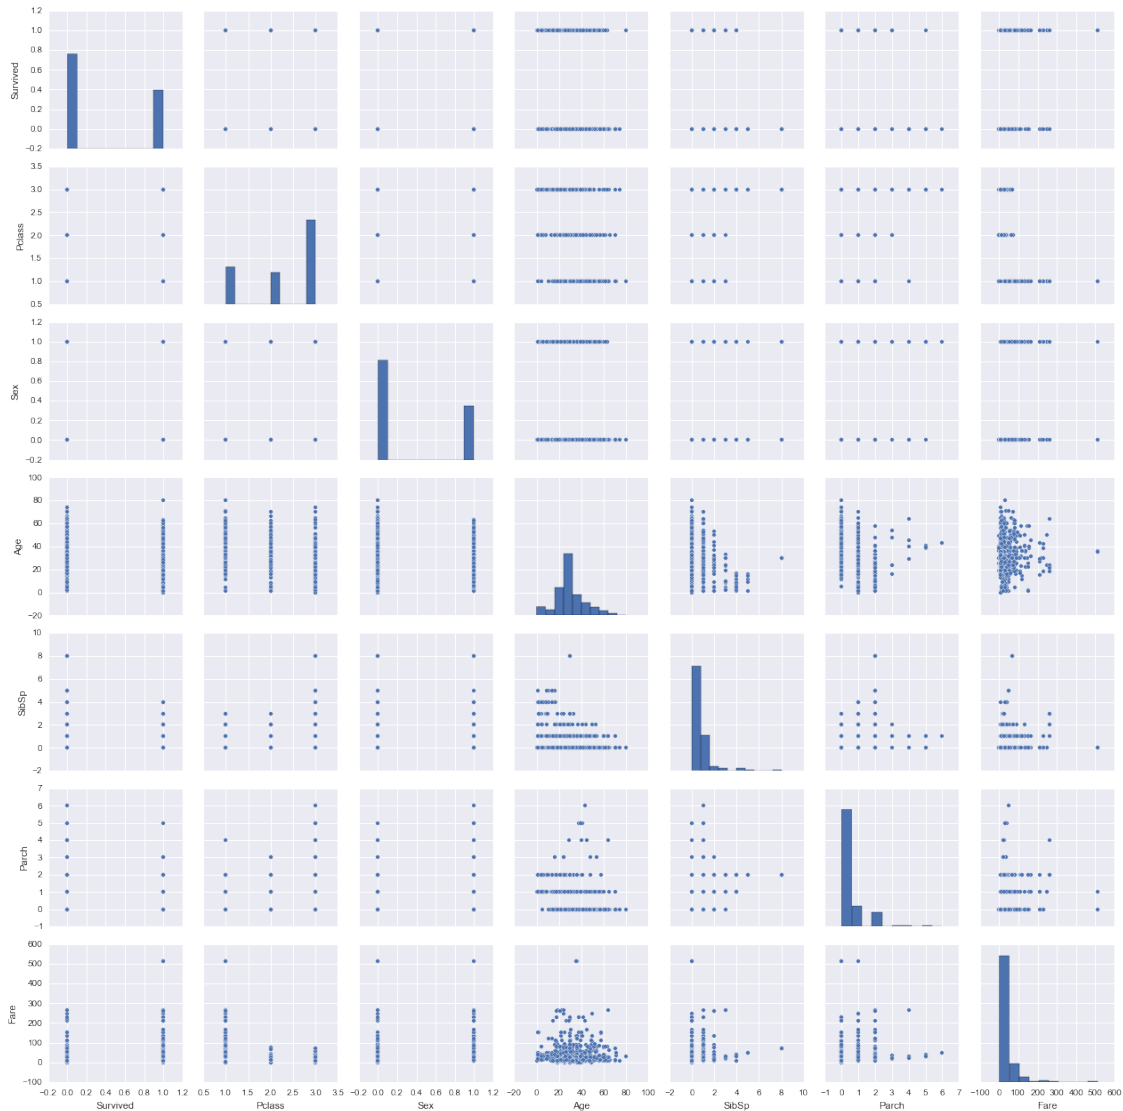


It appears are variables are not strongly correlated to each other. From the list of how the variables are correlated to the Survived label we can get the feeling a model might work quite well.

As a last step before we start training a model we will plot the pairwise distribution between the variables. Note: This only works because we only have a few columns. With more columns the figures look terrible.

```
In [10]: sns.pairplot(train_data)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x931aac8>
```



In [11]: *# Train the model*

```
train_x = train_data.drop('Survived', axis=1, inplace=False)
train_y = train_data.Survived
```

```
classifier = tree.DecisionTreeClassifier()
classifier.fit(train_x, train_y)
```

Out[11]: DecisionTreeClassifier(class\_weight=None, criterion='gini', max\_depth=None, max\_features=None, max\_leaf\_nodes=None, min\_samples\_leaf=1, min\_samples\_split=2, min\_weight\_fraction\_leaf=0.0, random\_state=None, splitter='best')

In [26]: *# Check score on training set*

```
cm_train = confusion_matrix(train_y, classifier.predict(train_x))
```

```

auc_train = roc_auc_score(train_y, classifier.predict(train_x))
f1_train = f1_score(train_y, classifier.predict(train_x))
print(cm_train)
print('-----')
print('AUC      : {}'.format(auc_train))
print('F1 Score : {}'.format(f1_train))

```

```

[[546  3]
 [ 16 326]]

```

```

-----
AUC      : 0.9738759466973445
F1 Score : 0.9716840536512668

```

```

In [15]: # Create predictions on the test set
         predictions = classifier.predict(test_data)

```

```

In [27]: # Score on the test outcomes
         cm_test = confusion_matrix(test_y, predictions)
         f1_test = f1_score(test_y, predictions)
         auc_test = roc_auc_score(test_y, predictions)
         print(cm_test)
         print('-----')
         print('AUC      : {}'.format(auc_test))
         print('F1 Score : {}'.format(f1_test))

```

```

[[225  41]
 [ 35 117]]

```

```

-----
AUC      : 0.8078007518796992
F1 Score : 0.7548387096774193

```

```

In [19]: # Possibility to save the tree to file, load it, turn it into PNG and show it here. D
         # Save decision tree to file, reload it and show in a plot
         #with open("dt.dot", 'w') as f:
         #    export_graphviz(classifier, out_file=f,
         #                    feature_names=list(train_x))
         #
         #    pydot.graph_from_dot_data(dotfile.getvalue()).write_png(file_path)
         #    i = misc.imread(file_path)
         #    plt.imshow(i)

```

There seems to be a lot of overfitting. Perhaps we can tweak the model to have some more constraints.

```

In [20]: # Look at our current tree
         classifier.get_params

```



```
Out [20]: <bound method DecisionTreeClassifier.get_params of DecisionTreeClassifier(class_weight=None,
max_features=None, max_leaf_nodes=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
random_state=None, splitter='best')>
```

```
In [21]: c2 = tree.DecisionTreeClassifier(max_depth = 15, max_leaf_nodes=10)
c2.fit(train_x, train_y)

f1_score(train_y, c2.predict(train_x))
```

```
Out [21]: 0.76609105180533754
```

We now have a highly constrained decision tree. A much lower score on the train set means we're not fitting to all our training examples anymore.  
Let's see if this helps our test predictions any

```
In [28]: pred = c2.predict(test_data)
print(confusion_matrix(test_y, pred))
print("-----")
print("AUC Score = {}".format(roc_auc_score(test_y, pred)))
print("F1 Score = {}".format(f1_score(test_y, pred)))
```

```
[[260   6]
 [ 12 140]]
```

```
-----
AUC Score = 0.9492481203007519
F1 Score  = 0.9395973154362416
```

We get a much better score with our constrained tree. Overfitting indeed was the problem.  
Finally let's try a Random Forest to see if a whole set of (constrained) decision trees does better than a single tree.

```
In [29]: forest = RandomForestClassifier(n_estimators=100, max_leaf_nodes=10)
forest.fit(train_x, train_y)
y_pred = forest.predict(test_data)
print("Accuracy score = {}".format(forest.score(test_data, test_y)))
print("F1 score      = {}".format(f1_score(test_y, y_pred)))
print("AUC score      = {}".format(roc_auc_score(test_y, y_pred)))
print("-----")
print(confusion_matrix(test_y, y_pred))
```

```
Accuracy score = 0.930622009569378
F1 score      = 0.9010238907849829
AUC score      = 0.9172932330827068
```

```
-----
[[257   9]
 [ 20 132]]
```